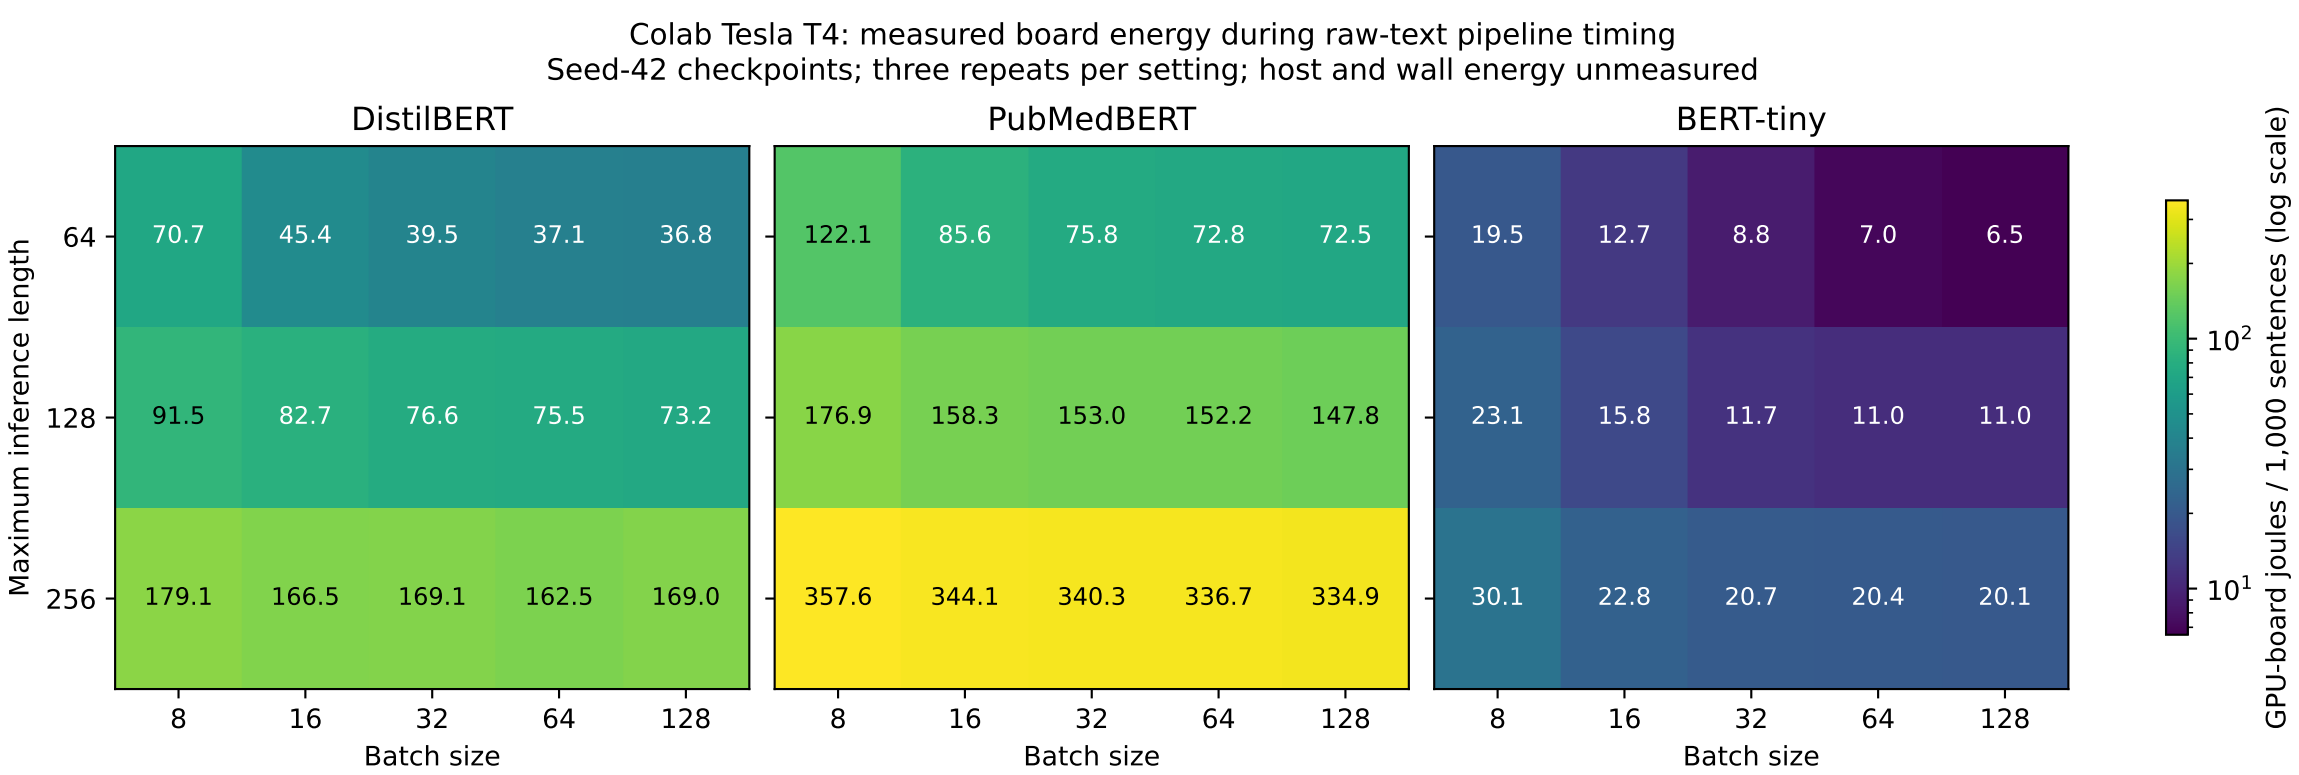
Colab Tesla T4: measured board energy during raw-text pipeline timing
Seed-42 checkpoints; three repeats per setting; host and wall energy unmeasured
DistilBERT
PubMedBERT
BERT-tiny
Maximum inference length
64
128
256
70.7
45.4
39.5
37.1
36.8
91.5
82.7
76.6
75.5
73.2
179.1
166.5
169.1
162.5
169.0
122.1
85.6
75.8
72.8
72.5
176.9
158.3
153.0
152.2
147.8
357.6
344.1
340.3
336.7
334.9
19.5
12.7
8.8
7.0
6.5
23.1
15.8
11.7
11.0
11.0
30.1
22.8
20.7
20.4
20.1
8
16
32
64
128
Batch size
GPU-board joules / 1,000 sentences (log scale)
10^2
10^1